

Google Reviews Pipeline for POIs (Aveiro)

Goal: Build end-to-end NLP pipeline to fetch, clean, and analyze reviews for Points of Interest (POIs) in Aveiro, using:

- **OSM-derived POIs** (from `pois_aveiro.csv` with EWKB geometry)
- **Google Places API (v1)** for proximity search and review extraction
- **Bilingual text preprocessing** (English & Portuguese with language-aware normalization)
- **NLP feature extraction** (TF-IDF for keywords, VADER sentiment for English)
- **Interactive visualizations** (ratings, sentiment, top places, wordclouds)
- Output **three-stage CSV pipeline**: raw → clean → enriched

Note: Google Places reviews are typically limited to 5 reviews per place.
Deduplication by `place_id` prevents duplicate results from overlapping search.

0. Setup & Requirements

Prerequisites

1. A **Google Cloud Project** with **Places API** enabled
2. A valid **API key** (set via environment variable `GOOGLE_API_KEY`)
3. The input file `pois_aveiro.csv` with `geom_pt` or `geom` column (EWKB POINT geometry, SRID=4326)

Key Features

- **Geometry parsing** from EWKB hex to (lon, lat) coordinates
- **Deduplication** across overlapping POI search circles
- **Bilingual NLP** (EN lemmatization + PT stemming with stopword removal)
- **Sentiment analysis** for English reviews (VADER)
- **TF-IDF keyword extraction** on processed text (without stopwords)
- **Multiple output formats** for exploratory analysis

Environment Variables

- `GOOGLE_API_KEY` : Your Google Places API key (fallback hardcoded for demo)

Rate Limits & Quotas

- Request delays are included to avoid hitting rate limits
- Reviews are limited to ~5 per place (Google API constraint)
- Adjust `MAX_TOTAL_POIS` to control overall execution time

```
In [ ]: import os      # filesystem paths, env vars
INPUT_CSV_PRIMARY = "../../Milestone_2/pois_aveiro.csv"
INPUT_CSV_FALLBACK = "../../Milestone_2/pois_aveiro.csv"
OUTPUT_DIR = "../output"
RAW_REVIEWS_CSV = os.path.join(OUTPUT_DIR, "reviews_raw.csv")
CLEAN_REVIEWS_CSV = os.path.join(OUTPUT_DIR, "reviews_clean.csv")

# Google Places API
GOOGLE_API_KEY = os.getenv("GOOGLE_API_KEY", "not_a_real_key")

# Fetch configuration
SLEEP_BETWEEN_REQUESTS = 0.1    # seconds
RADIUS_METERS = 25              # tighter radius to match OSM POI
MAX_PLACES_PER_POI = 1000       # max nearby results per POI we will consider
MAX_TOTAL_POIS = 1000           # safety limit for demonstration; set None for all

INCLUDED_TYPES = None            # e.g., ["restaurant", "cafe"] or None for all

# Retry configuration
MAX_RETRIES = 3
BACKOFF_FACTOR = 1.6

os.makedirs(OUTPUT_DIR, exist_ok=True)
print("Output directory:", OUTPUT_DIR)
```

Output directory: ../output

```
In [ ]: # === Core & Utilities ===
import time      # simple delays between requests
import json      # summary output
import math      # any minor math utilities
import re         # regex for text cleaning
import ast        # safe literal evaluation for parsed dicts

# === HTTP & DataFrames ===
import requests   # Google Places API calls
import pandas as pd # data manipulation

# === Geometry ===
from shapely import wkb # EWKB hex POINT parsing

# === Typing helpers ===
from typing import Optional, Tuple, List, Dict, Any

# === NLP & Language ===
import numpy as np                         # numeric ops
import nltk                                # tokenization/resources
from nltk.corpus import stopwords
from nltk.stem import WordNetLemmatizer, RSLPStemmer
from nltk.tokenize import word_tokenize
try:
    from langdetect import detect             # language detection (optional)
    HAVE_LANGDETECT = True
except Exception:
    HAVE_LANGDETECT = False

# === Features & Sentiment ===
from sklearn.feature_extraction.text import TfidfVectorizer
try:
    from vaderSentiment.vaderSentiment import SentimentIntensityAnalyzer
```

```
HAVE_VADER = True
    sid = SentimentIntensityAnalyzer()
except Exception:
    HAVE_VADER = False
    sid = None

# === Visualization ===
import matplotlib.pyplot as plt
import seaborn as sns
try:
    from wordcloud import WordCloud
    HAVE_WORDCLOUD = True
except Exception:
    HAVE_WORDCLOUD = False

# === Transformers (optional for PT sentiment) ===
try:
    from transformers import pipeline
    HAVE_TRANSFORMERS = True
except Exception:
    HAVE_TRANSFORMERS = False

# === Topic Modeling (optional for LDA) ===
try:
    from gensim import corpora
    from gensim.models import LdaMulticore
    HAVE_GENSIM = True
except ImportError:
    HAVE_GENSIM = False

print("Imports loaded. Optional libs:", {
    "langdetect": HAVE_LANGDETECT,
    "vader": HAVE_VADER,
    "wordcloud": HAVE_WORDCLOUD,
    "transformers": HAVE_TRANSFORMERS,
    "gensim": HAVE_GENSIM,
})
```

```
/home/miragaia/Documents/5_ANO/ICD/PROJECT_ICD/.venv/lib/python3.12/site-packages/tqdm/auto.py:21: TqdmWarning: IProgress not found. Please update jupyter and ipywidgets. See https://ipywidgets.readthedocs.io/en/stable/user_install.html
```

```
from .autonotebook import tqdm as notebook_tqdm
Imports loaded. Optional libs: {'langdetect': True, 'vader': True, 'wordcloud': True, 'transformers': True}
```

1. Geometry Parsing & API Client Functions

Purpose: Define helper functions for geometric transformations and Google Places API interactions.

Process

1. **EWKB to Coordinates:** Parse OSM EWKB hex-encoded POINT geometries (SRID 4326) into standard (lon, lat) tuples
2. **API Headers & Field Masking:** Construct proper HTTP headers with field masks to request specific place attributes (name, ID, reviews, location, rating, type)

3. **Nearby Search:** Query Google Places API for places within a search radius around each POI
4. **Review Extraction:** Parse the API response and flatten nested review structures into rows

Key Functions

- `ewkb_hex_point_to_lonlat()` : Convert EWKB hex to geographic coordinates
- `places_search_nearby()` : Execute Nearby Search query (single page, no pagination)
- `extract_reviews_from_places()` : Extract and flatten reviews from place objects

Notes

- EWKB format commonly appears in PostGIS/OSM exports (e.g.,
0101000020E6100000...)
- Single-page Nearby Search avoids field mask conflicts; pagination available as future enhancement
- Deduplication of places by ID prevents duplicate reviews from overlapping search circles

```
In [3]: def ewkb_hex_point_to_lonlat(hex_str: str) -> Optional[Tuple[float, float]]:
    """
    Convert EWKB hex POINT (SRID=4326) to (lon, lat).
    """
    if not isinstance(hex_str, str) or not hex_str:
        return None
    try:
        geom = wkb.loads(bytes.fromhex(hex_str))
        if geom.geom_type == "Point":
            return (geom.x, geom.y)
    except Exception:
        pass
    return None

def _places_headers(field_mask: str) -> Dict[str, str]:
    return {
        "X-Goog-FieldMask": field_mask,
        "X-Goog-Api-Key": GOOGLE_API_KEY,
        "Content-Type": "application/json",
    }

def places_search_nearby(
    lat: float,
    lon: float,
    radius: int = 25,
    included_types: Optional[List[str]] = None,
    max_results: int = 20,
) -> List[Dict[str, Any]]:
    """
    Nearby Search (single page) following the Milestone 2 approach.
    Returns a list of place dicts (no pagination).
    """
```

```
"""
url = "https://places.googleapis.com/v1/places:searchNearby"

payload: Dict[str, Any] = {
    "locationRestriction": {
        "circle": {
            "center": {"latitude": lat, "longitude": lon},
            "radius": radius,
        }
    },
    "maxResultCount": min(max_results, 20),
}
if included_types:
    payload["includedTypes"] = included_types

field_mask = (
    "places.displayName,places.id,places.reviews,places.location,pla
)

try:
    response = requests.post(
        url,
        params={"key": GOOGLE_API_KEY},
        json=payload,
        headers=_places_headers(field_mask),
        timeout=20,
    )
    response.raise_for_status()
except requests.exceptions.RequestException as e:
    print(f"Error during API request: {e}")
    return []

data = response.json()
places = data.get("places", []) if isinstance(data, dict) else []
return places


def extract_reviews_from_places(places: List[Dict[str, Any]]) -> List[Dic
rows = []
for p in places:
    pid = p.get("id")
    pname = (p.get("displayName") or {}).get("text", None)
    prating = p.get("rating")
    ptype = p.get("primaryType")
    ploc = p.get("location") or {}
    reviews = p.get("reviews") or []
    for r in reviews:
        rows.append({
            "place_id": pid,
            "place_name": pname,
            "place_rating": prating,
            "place_primary_type": ptype,
            "place_location": ploc,
            "author_name": (r.get("authorAttribution") or {}).get("di
            "rating": r.get("rating"),
            "review_text": (r.get("text") or {}).get("text", ""),
            "publish_time": r.get("publishTime"),
        })
return rows
```

2. Load POIs & Fetch Reviews via Nearby Search

Purpose: Iterate over OSM-derived POIs, execute Nearby Search for each location, and aggregate all discovered reviews into a **raw output CSV**.

⚠ SKIP THIS SECTION if `reviews_raw.csv` is available and want to work with existing data (see Section 2b instead).

Process

1. **Load POI CSV:** Read `pois_aveiro.csv` and verify it contains EWKB geometry column (`geom_pt` or `geom`)
2. **Parse Coordinates:** For each POI row, extract (lon, lat) from EWKB hex string
3. **Execute Nearby Search:** Query Google Places API within `RADIUS_METERS` around each POI to discover nearby places
4. **Deduplicate Places:** Track `place_id` globally to prevent duplicate processing from overlapping search circles
5. **Extract Reviews:** For each unique place, flatten nested review structures into individual rows
6. **Attach Context:** Enrich each review with POI metadata (gid, amenity type, source name)
7. **Parse Location:** Convert embedded place location objects to separate `lat` and `lon` columns for convenience
8. **Save Raw Output:** Write all reviews to `reviews_raw.csv` with original fields + POI context

Key Parameters

- `RADIUS_METERS` : Search radius around each POI (25m = roughly OSM POI precision)
- `MAX_PLACES_PER_POI` : Limit results per Nearby Search query (1000 = API max)
- `MAX_TOTAL_POIS` : Safety limit for total POIs processed (set `None` to process all)
- `INCLUDED_TYPES` : Optional list to filter places (e.g., `["restaurant", "cafe"]`), or `None` for all types

Output

- `reviews_raw.csv` : Raw API results with columns: `place_id`, `place_name`, `author_name`, `rating`, `review_text`, `publish_time`, `lat`, `lon`, `place_rating`, `place_primary_type`, `poi_context`...
- Deduplication count printed to show unique places discovered

Alternative: If you already have fetched reviews and lost API access, skip to **Section 2b** to load from existing CSV.

```
In [5]: if os.path.exists(INPUT_CSV_PRIMARY):  
    pois_path = INPUT_CSV_PRIMARY
```

```
elif os.path.exists(INPUT_CSV_FALLBACK):
    pois_path = INPUT_CSV_FALLBACK
else:
    raise FileNotFoundError(
        f"Could not find input CSV at {INPUT_CSV_PRIMARY} or fallback {IN
    )

print("Loading:", pois_path)
df_pois = pd.read_csv(pois_path, low_memory=False)
print("Rows:", len(df_pois))

geom_col = None
for cand in ["geom_pt", "geom"]:
    if cand in df_pois.columns:
        geom_col = cand
        break
if geom_col is None:
    raise ValueError("No EWKB point column ('geom_pt' or 'geom') found in

out_rows = []
processed = 0
seen_place_ids = set()

for idx, row in df_pois.iterrows():
    if MAX_TOTAL_POIS and processed >= MAX_TOTAL_POIS:
        break

    hex_point = row.get(geom_col)
    lonlat = ewkb_hex_point_to_lonlat(hex_point) if isinstance(hex_point,
    if not lonlat:
        continue

    lon, lat = lonlat

    places = places_search_nearby(
        lat=lat,
        lon=lon,
        radius=RADIUS_METERS,
        included_types=INCLUDED_TYPES,
        max_results=MAX_PLACES_PER_POI,
    )

    unique_places = []
    for p in places:
        pid = p.get("id")
        if pid and pid not in seen_place_ids:
            unique_places.append(p)
            seen_place_ids.add(pid)

    rows = extract_reviews_from_places(unique_places)

    for r in rows:
        r["poi_row_index"] = idx
        if "gid" in df_pois.columns:
            r["poi_gid"] = row.get("gid")
        for c in ["amenity", "shop", "tourism", "name"]:
            if c in df_pois.columns:
                r[f"poi_{c}"] = row.get(c)

    out_rows.extend(rows)
```

```
processed += 1

if out_rows:
    df_raw = pd.DataFrame(out_rows)
def parse_loc(x):
    if isinstance(x, dict):
        return x
    try:
        return ast.literal_eval(x)
    except Exception:
        return {}

df_raw["place_location"] = df_raw["place_location"].apply(parse_loc)
df_raw["lat"] = df_raw["place_location"].apply(lambda d: d.get("lat"))
df_raw["lon"] = df_raw["place_location"].apply(lambda d: d.get("longi"))

df_raw.to_csv(RAW_REVIEWS_CSV, index=False)
print(f"Saved {len(df_raw)} raw reviews -> {RAW_REVIEWS_CSV}")
print(f"Total unique places fetched: {len(seen_place_ids)}")
else:
    df_raw = pd.DataFrame()
    print("No reviews fetched.")

display(df_raw.head(10))
```

```
Loading: ../../Milestone_2/pois_aveiro.csv
Rows: 13258
Rows: 13258
Error during API request: 500 Server Error: Internal Server Error for url:
https://places.googleapis.com/v1/places:searchNearby?key=AIzaSyAwme_k4xStLv2_bLFAckn55hJ1TNF8L8A
Error during API request: 500 Server Error: Internal Server Error for url:
https://places.googleapis.com/v1/places:searchNearby?key=AIzaSyAwme_k4xStLv2_bLFAckn55hJ1TNF8L8A
Saved 6800 raw reviews -> ../output/reviews_raw.csv
Total unique places fetched: 2356
Saved 6800 raw reviews -> ../output/reviews_raw.csv
Total unique places fetched: 2356
```

	place_id	place_name	place_rating	place_primary_type	
0	ChIJgymEZASYlw0RKRhKH9TfRQ	bp	4.1	gas_station	40.63
1	ChIJgymEZASYlw0RKRhKH9TfRQ	bp	4.1	gas_station	40.63
2	ChIJgymEZASYlw0RKRhKH9TfRQ	bp	4.1	gas_station	40.63
3	ChIJgymEZASYlw0RKRhKH9TfRQ	bp	4.1	gas_station	40.63
4	ChIJgymEZASYlw0RKRhKH9TfRQ	bp	4.1	gas_station	40.63
5	ChIJdZmB9g-jlw0RBVGjwMFX79Y	Café BP Maia	4.3	cafe	'longi
6	ChIJdZmB9g-jlw0RBVGjwMFX79Y	Café BP Maia	4.3	cafe	'longi
7	ChIJdZmB9g-jlw0RBVGjwMFX79Y	Café BP Maia	4.3	cafe	'longi
8	ChIJdZmB9g-jlw0RBVGjwMFX79Y	Café BP Maia	4.3	cafe	'longi
9	ChIJdZmB9g-jlw0RBVGjwMFX79Y	Café BP Maia	4.3	cafe	'longi

2b. Load Existing Reviews from CSV (Skip API Fetch)

Purpose: Load previously fetched reviews from `reviews_raw.csv` when API access is unavailable or when working with existing data.

Use Case

- **API quota exhausted** or API key no longer available
- **Working offline** with previously collected data
- **Reproducible analysis** using a fixed snapshot of reviews

- **Faster iteration** during development (skip expensive API calls)

Process

1. **Check for existing CSV:** Verify `reviews_raw.csv` exists in the output directory
2. **Load DataFrame:** Read CSV with all original columns (`place_id`, `place_name`, `author_name`, `rating`, `review_text`, etc.)
3. **Parse location data:** Convert `place_location` from string representation back to dictionary if needed
4. **Extract coordinates:** Ensure `lat` and `lon` columns are properly populated
5. **Validate data:** Check for required columns and print summary statistics

Key Columns Expected

- `place_id` : Google Place ID (unique identifier)
- `place_name` : Name of the place
- `author_name` : Review author
- `rating` : Star rating (1-5)
- `review_text` : Original review text
- `publish_time` : ISO timestamp of review
- `lat` , `lon` : Geographic coordinates
- `place_rating` : Overall place rating
- `place_primary_type` : Place category (restaurant, cafe, etc.)
- `poi_*` columns: Original POI metadata from OSM

Output

- `df_raw` : Pandas DataFrame ready for downstream processing (identical structure to API fetch output)
- Summary statistics printed (total reviews, date range, language distribution preview)

Note: Run **either** Section 2 (API fetch) **OR** Section 2b (load from CSV), not both.
Comment out the section you don't need.

```
In [ ]: # Load existing reviews from CSV (alternative to API fetch in Section 2)
if os.path.exists(RAW_REVIEWS_CSV):
    print(f"Loading existing reviews from: {RAW_REVIEWS_CSV}")
    df_raw = pd.read_csv(RAW_REVIEWS_CSV, low_memory=False)

# Parse place_location if it's stored as string
def parse_loc(x):
    if isinstance(x, dict):
        return x
    if pd.isna(x):
        return {}
    try:
        return ast.literal_eval(str(x))
    except Exception:
        return {}

# Ensure location parsing and coordinate extraction
```

```
if "place_location" in df_raw.columns:
    df_raw["place_location"] = df_raw["place_location"].apply(parse_l

# Extract or verify lat/lon columns
if "lat" not in df_raw.columns or df_raw["lat"].isna().all():
    df_raw["lat"] = df_raw["place_location"].apply(lambda d: d.get("l
if "lon" not in df_raw.columns or df_raw["lon"].isna().all():
    df_raw["lon"] = df_raw["place_location"].apply(lambda d: d.get("l

print(f"Loaded {len(df_raw)} reviews")
print(f"Unique places: {df_raw['place_id'].nunique()}")
print(f"Date range: {df_raw['publish_time'].min()} to {df_raw['publis
print(f"Average rating: {df_raw['rating'].mean():.2f}")

display(df_raw.head(10))
else:
    raise FileNotFoundError(
        f"CSV file not found: {RAW_REVIEWS_CSV}\n"
        f"Please run Section 2 (API fetch) first to generate the raw revi
        f"or check that the OUTPUT_DIR path is correct."
    )
```

Loading existing reviews from: ../output/reviews_raw.csv
Loaded 6800 reviews
Unique places: 1655
Date range: 2009-10-22T19:26:04.476521Z to 2025-12-08T08:31:55.750368399Z
Average rating: 4.32

	place_id	place_name	place_rating	place_primary_type	
0	ChIJgymEZASYlw0RKRhKH9TfRQ	bp	4.1	gas_station	40.63
1	ChIJgymEZASYlw0RKRhKH9TfRQ	bp	4.1	gas_station	40.63
2	ChIJgymEZASYlw0RKRhKH9TfRQ	bp	4.1	gas_station	40.63
3	ChIJgymEZASYlw0RKRhKH9TfRQ	bp	4.1	gas_station	40.63
4	ChIJgymEZASYlw0RKRhKH9TfRQ	bp	4.1	gas_station	40.63
5	ChIJdZmB9g-jlw0RBVGjwMFX79Y	Café BP Maia	4.3	cafe	'longi
6	ChIJdZmB9g-jlw0RBVGjwMFX79Y	Café BP Maia	4.3	cafe	'longi
7	ChIJdZmB9g-jlw0RBVGjwMFX79Y	Café BP Maia	4.3	cafe	'longi
8	ChIJdZmB9g-jlw0RBVGjwMFX79Y	Café BP Maia	4.3	cafe	'longi
9	ChIJdZmB9g-jlw0RBVGjwMFX79Y	Café BP Maia	4.3	cafe	'longi

3. Bilingual Text Preprocessing & Normalization

Purpose: Clean raw review text and normalize it using language-aware techniques, producing a `text_processed` column without stopwords for downstream NLP tasks.

Input: Uses `df_raw` DataFrame from either Section 2 (API fetch) or Section 2b (load from CSV)

Process

- 1. Text Cleaning:** Remove HTML tags, URLs, excessive whitespace, and non-

alphanumeric characters (preserving accented characters for Portuguese/Spanish)

2. **Language Detection:** Use `langdetect` to identify whether each review is in English (en), Portuguese (pt), or another language; optionally filter to EN/PT only

3. **Language-Specific Tokenization:**

- **English:** Tokenize → remove EN stopwords → lemmatize (WordNet) → keep tokens > 2 chars
- **Portuguese:** Tokenize → remove PT stopwords → stem (RSLP - Snowball Stemmer for Portuguese) → keep tokens > 2 chars
- **Other languages:** Tokenize → remove bilingual stopwords → apply English lemmatization as fallback

4. **Reconstruct Processed Text:** Join tokens back into `text_processed` column (used by TF-IDF and wordcloud)

5. **Compute Token Stats:** Count tokens per review for exploratory analysis

6. **Save Clean Output:** Write to `reviews_clean.csv` with added columns: `review_text_clean`, `lang`, `tokens`, `token_count`, `text_processed`

Language-Specific Methods

- **English Lemmatization:** WordNet lemmatizer (reduces "running", "runs" → "run")
- **Portuguese Stemming:** RSLP stemmer (reduces "casas", "casa" → "cas")
- **Stopword Removal:** Pre-computed sets for EN (179 words) and PT (228 words)

Key Columns in Output

- `review_text_clean` : Original text with only punctuation/URLs removed (light cleaning)
- `lang` : Detected language code (en, pt, or null if detection failed)
- `tokens` : List of normalized tokens (stopwords already removed)
- `text_processed` : String reconstructed from tokens (ready for TF-IDF and visualization)

Graceful Fallbacks

- If `langdetect` unavailable: skips language detection, uses bilingual stopword list
- Unknown languages: defaults to English lemmatization
- Missing or empty reviews: handled as empty token lists

Notes

- `text_processed` is the key column for downstream TF-IDF and wordcloud (not `review_text_clean`)
- Stopword removal + lemmatization/stemming significantly reduces noise in keyword extraction

```
In [ ]: # Stopwords for EN and PT
stop_en = set(stopwords.words('english'))
try:
    stop_pt = set(stopwords.words('portuguese'))
except Exception:
```

```

stop_pt = set()
stop_all = stop_en.union(stop_pt)

lemmatizer = WordNetLemmatizer()
rslp = RSLPStemmer()

def clean_text(s: str) -> str:
    s = str(s)
    s = re.sub(r"<[^>]+>", " ", s) # HTML
    s = re.sub(r"http\S+|www\S+", " ", s) # URLs
    s = re.sub(r"[^\w\séíóúàèòâêôçõõÁÉÍÓÚÀÈÌÒÙÂÊÎÔÛÇÃÕ]", " ", s) #
    s = re.sub(r"\s+", " ", s)
    return s.strip().lower()

def language_or_none(s: str) -> str:
    if not HAVE_LANGDETECT:
        return None
    try:
        return detect(s)
    except Exception:
        return None

def simple_tokenize(s: str) -> list:
    # Use nltk's word_tokenize; fallback to regex if it fails
    try:
        return [t for t in word_tokenize(s) if t]
    except Exception:
        return.findall(r"\b\w+\b", s, flags=re.UNICODE)

def tokens_for_lang(s: str, lang: str | None) -> list:
    toks = simple_tokenize(s)
    if lang == 'pt':
        toks = [t for t in toks if t.isalpha() and t not in stop_pt and l
        toks = [rslp.stem(t) for t in toks]
        return toks
    elif lang == 'en':
        toks = [t for t in toks if t.isalpha() and t not in stop_en and l
        toks = [lemmatizer.lemmatize(t) for t in toks]
        return toks
    else:
        # Unknown language: use a bilingual stopword list, no language-specific
        toks = [t for t in toks if t.isalpha() and t.lower() not in stop_
        # Default to English lemmatizer as a light normalization
        toks = [lemmatizer.lemmatize(t) for t in toks]
        return toks

if not df_raw.empty:
    df_clean = df_raw.copy()
    df_clean["review_text_clean"] = df_clean["review_text"].fillna("").ma

    # Language detection (EN/PT only)
    if HAVE_LANGDETECT:
        df_clean["lang"] = df_clean["review_text_clean"].map(language_or_
        df_clean = df_clean[df_clean["lang"].isin(["en", "pt"])].fillna(Tr
    else:

```

```
df_clean["lang"] = np.nan

# Tokenization and normalization per language
df_clean["tokens"] = df_clean.apply(lambda r: tokens_for_lang(r["review"]),
df_clean["token_count"] = df_clean["tokens"].map(len)
# Reconstruct a processed text without stopwords for downstream features
df_clean["text_processed"] = df_clean["tokens"].map(lambda toks: " ".join(toks))

df_clean.to_csv(CLEAN_REVIEWS_CSV, index=False)
print(f"Saved clean reviews -> {CLEAN_REVIEWS_CSV}")

else:
    df_clean = pd.DataFrame()

display(df_clean.head(10))
```

Saved clean reviews -> ../output/reviews_clean.csv

	place_id	place_name	place_rating	place_primary_type
0	ChIJgymEZASYlw0RKRhKH9TfRQ	bp	4.1	gas_station 40.63
1	ChIJgymEZASYlw0RKRhKH9TfRQ	bp	4.1	gas_station 40.63
2	ChIJgymEZASYlw0RKRhKH9TfRQ	bp	4.1	gas_station 40.63
3	ChIJgymEZASYlw0RKRhKH9TfRQ	bp	4.1	gas_station 40.63
4	ChIJgymEZASYlw0RKRhKH9TfRQ	bp	4.1	gas_station 40.63
5	ChIJdZmB9g-jlw0RBVGjwMFX79Y	Café BP Maia	4.3	cafe 'longi
6	ChIJdZmB9g-jlw0RBVGjwMFX79Y	Café BP Maia	4.3	cafe 'longi
7	ChIJdZmB9g-jlw0RBVGjwMFX79Y	Café BP Maia	4.3	cafe 'longi
8	ChIJdZmB9g-jlw0RBVGjwMFX79Y	Café BP Maia	4.3	cafe 'longi
9	ChIJdZmB9g-jlw0RBVGjwMFX79Y	Café BP Maia	4.3	cafe 'longi

10 rows × 22 columns

4. Feature Extraction & Sentiment Analysis

Purpose: Compute quantitative features from processed reviews (keyword importance via TF-IDF) and bilingual sentiment polarity scores.

Process

1. **TF-IDF Vectorization:** Fit TF-IDF model on `text_processed` (stopword-removed, normalized text)

- Extract unigrams and bigrams (1-2 word phrases)
- Minimum document frequency = 2 (appears in ≥2 reviews to reduce noise)
- Maximum features = 5000 (most frequent/important terms)
- Aggregate to compute mean TF-IDF per place

2. **Bilingual Sentiment Analysis**

- **English:** Apply VADER (Valence Aware Dictionary and sEntiment Reasoner) to `review_text_clean`
 - VADER designed for social media text; returns compound score in [-1, 1] range
 - Score = -1 (most negative), 0 (neutral), +1 (most positive)
- **Portuguese:** Use HuggingFace `ysentimiento/bertweet-pt-sentiment` (text classification)
 - We compute a continuous compound score as:
$$\text{compound} = P(\text{POSITIVE}) - P(\text{NEGATIVE})$$
, using the model's softmax probabilities
 - This yields a value in approximately [-1, 1]; NEUTRAL mass is implicit
 - If Transformers are unavailable, fallback to TextBlob polarity in [-1, 1]
 - Both methods preserve punctuation and text emphasis

3. **Per-Place Keyword Extraction:** For each `place_id`, compute top 10 keywords by mean TF-IDF score

4. **Save Enriched Output:** Write to `reviews_enriched.csv` with added columns: `tf_idf_features`, `sentiment_compound` (bilingual)

Key Outputs

- `top_keywords_per_place` : Dictionary mapping `place_id` → [(keyword, tfidf_score), ...]
- `sentiment_compound` : Float in [-1, 1] for both EN (VADER) and PT (HuggingFace or TextBlob fallback)
- `X` : Sparse TF-IDF matrix (scipy.sparse CSR format) for downstream modeling

Notes

- Use `text_processed` for TF-IDF (no stopwords, normalized) to get meaningful keywords
- Use `review_text_clean` for sentiment (punctuation preserved, needed for VADER emphasis detection)
- Bilingual sentiment enables comparative analysis across languages with consistent scale
- `bertweet-pt` is stronger for Portuguese than simple lexicon-based approaches
- Both TF-IDF and sentiment results saved to enable further analysis (topic

modeling, emotion classification, etc.)

```
In [6]: if not df_clean.empty:
    # Reset index so TF-IDF matrix rows align with DataFrame rows
    df_idx = df_clean.reset_index(drop=True)

    # TF-IDF on processed text (stopwords removed, lemmatized/stemmed)
    texts = df_idx["text_processed"].fillna("").tolist()
    tfidf = TfidfVectorizer(max_features=5000, ngram_range=(1,2), min_df=X)
    X = tfidf.fit_transform(texts)
    vocab = np.array(tfidf.get_feature_names_out())

    # Compute per-place top keywords by mean TF-IDF
    top_keywords_per_place = {}
    for pid, sub in df_idx.groupby("place_id"):
        idx = sub.index
        if len(idx) == 0:
            continue
        vec = X[idx].mean(axis=0)
        arr = np.asarray(vec).ravel()
        top_idx = arr.argsort()[-10:][::-1]
        top_keywords_per_place[pid] = list(zip(vocab[top_idx], arr[top_idx]))

    # Sentiment: VADER for English, HuggingFace (bertweet-pt) for Portuguese
    df_clean["sentiment_compound"] = np.nan

    # English sentiment (VADER)
    if HAVE_VADER:
        en_mask = df_clean["lang"].eq("en")
        df_clean.loc[en_mask, "sentiment_compound"] = df_clean[en_mask][
            lambda s: sid.polarity_scores(s)["compound"]
        )

    # Portuguese sentiment (HuggingFace bertweet-pt-sentiment). Fallback
    pt_mask = df_clean["lang"].eq("pt")
    pt_senti = None
    if HAVE_TRANSFORMERS:
        try:
            pt_senti = pipeline("text-classification", model="pysentimentien")
        except Exception as e:
            print("Transformers pipeline init failed:", e)
            pt_senti = None

    def pt_compound_score(text: str) -> float:
        if not isinstance(text, str) or not text.strip():
            return np.nan
        if pt_senti is None:
            return np.nan
        try:
            # Use top_k=None to get all class scores (replaces deprecated
            res = pt_senti(text, top_k=None)
            # Expected format: [{"label": 'NEG', 'score': 0.x}, {"label": ...
            if isinstance(res, list) and len(res) > 0:
                # Build score dict from all returned labels
                scores = {}
                for item in res:
                    label = str(item.get("label", "")).upper()
                    score_val = float(item.get("score", 0.0))
                    # Map various label formats to standard names
                    if label in ["NEG", "NEU", "POS"]:
                        scores[label] = score_val
                    else:
                        scores["NEU"] = score_val
```
```

```
 if "POS" in label:
 scores["POSITIVE"] = score_val
 elif "NEG" in label:
 scores["NEGATIVE"] = score_val
 elif "NEU" in label:
 scores["NEUTRAL"] = score_val

 ppos = scores.get("POSITIVE", 0.0)
 pneg = scores.get("NEGATIVE", 0.0)
 # Compute compound as difference: positive - negative (ra
 compound = float(ppos - pneg)
 return compound
return 0.0 # Fallback if format unexpected
except Exception as e:
 # Silent fallback for individual review errors
 return np.nan

if pt_senti is not None:
 df_clean.loc[pt_mask, "sentiment_compound"] = df_clean.loc[pt_mas
 print("Portuguese sentiment computed with bertweet-pt (HuggingFac
else:
 try:
 from textblob import TextBlob
 df_clean.loc[pt_mask, "sentiment_compound"] = df_clean.loc[pt_
 lambda s: TextBlob(s).sentiment.polarity
)
 print("Transformers unavailable; Portuguese sentiment via Tex
 except ImportError:
 print("TextBlob not available; Portuguese sentiment skipped.

Persist enriched
df_enriched = df_clean.copy()
df_enriched.to_csv(os.path.join(OUTPUT_DIR, "reviews_enriched.csv"),
print("Saved reviews_enriched.csv")
else:
 top_keywords_per_place = {}
 df_enriched = pd.DataFrame()

Preview keywords for first 3 places
for i, (pid, kws) in enumerate(top_keywords_per_place.items()):
 if i >= 3:
 break
 print(f"Place {pid} top keywords:")
 for w, score in kws:
 print(" ", w, f"{score:.3f}")
```

emoji is not installed, thus not converting emoticons or emojis into text.  
Install emoji: pip3 install emoji==0.6.0  
Device set to use cpu  
Device set to use cpu

Portuguese sentiment computed with bertweet-pt (HuggingFace)  
Saved reviews\_enriched.csv  
Place ChIJ-1iL6JiZIw0RDUrtC7a89ZI top keywords:  
girl 0.134  
happy 0.111  
cabel 0.109  
gabriel 0.106  
first 0.099  
desd 0.071  
result 0.070  
incre 0.065  
destaqu 0.065  
sent bem 0.061  
Place ChIJ-6D9iC0XIw0RrrI1s929jkA top keywords:  
cabel 0.294  
ter 0.105  
equip incr 0.103  
real 0.100  
faz 0.099  
trat 0.090  
poss 0.086  
result 0.083  
menin 0.080  
loir 0.077  
Place ChIJ-7D2DEyRIw0RajlTjh1ovls top keywords:  
sunset 0.131  
beach 0.116  
served 0.105  
food 0.088  
place right 0.079  
food drink 0.077  
crowded 0.077  
cocktail 0.074  
reccomend 0.073  
friendly 0.071

## 5. Exploratory Data Visualization (Bilingual)

**Purpose:** Generate visual summaries split by detected language (English vs. Portuguese) to compare patterns across linguistic contexts.

### Charts Produced

#### 1. Rating Distributions

- Separate histograms with KDE for English and Portuguese reviews
- Shows star rating (1-5) frequency distribution per language

#### 2. Sentiment Distributions (Bilingual)

- **English:** VADER compound scores [-1, 1] with histogram + KDE
- **Portuguese:** HuggingFace bertweet-pt compound scores [-1, 1] with histogram + KDE
- Enables direct comparison of sentiment polarity across languages

#### 3. Top Places by Review Count

- Bar charts showing top 10 most-reviewed places per language
- Identifies popular venues within each linguistic community

#### 4. Word Clouds (Language-Specific Tokenization)

- **English:** Lemmatized tokens from `text_processed`
- **Portuguese:** Unstemmed tokens (preserves readability) with stopword removal
- Visual representation of most frequent terms per language

#### Notes

- All visualizations gracefully skip empty language slices with informative messages
- Portuguese sentiment now uses the same [-1, 1] scale as English for direct comparison
- Word clouds use different preprocessing: EN uses lemmas, PT uses raw tokens for better interpretability

```
In [7]: # Bilingual visualizations and comparative analysis
if df_enriched.empty:
 print("No data to visualize.")
else:
 df_en = df_enriched[df_enriched["lang"].eq("en")]
 df_pt = df_enriched[df_enriched["lang"].eq("pt")]

 # Helper functions for visualizations
 def plot_ratings(df_lang, label):
 if df_lang.empty:
 print(f"No {label} reviews for rating histogram.")
 return
 plt.figure(figsize=(7,4))
 sns.histplot(pd.to_numeric(df_lang["rating"], errors="coerce").dropna(), bins=30, kde=True)
 plt.title(f"Rating Distribution ({label})")
 plt.xlabel("Rating")
 plt.ylabel("Count")
 plt.tight_layout(); plt.show()

 def plot_sentiment(df_lang, label, method_name):
 """Plot sentiment distribution for any language with proper label"""
 if df_lang.empty or df_lang["sentiment_compound"].notna().sum() == 0:
 print(f"No {label} sentiment data to plot.")
 return
 plt.figure(figsize=(7,4))
 sns.histplot(df_lang["sentiment_compound"].dropna(), bins=30, kde=True)
 plt.title(f"Sentiment Distribution ({label}, {method_name})")
 plt.xlabel("Compound Score [-1, 1]")
 plt.ylabel("Count")
 plt.tight_layout(); plt.show()

 def plot_top_places(df_lang, label):
 if df_lang.empty:
 print(f"No {label} reviews for top places.")
 return
 top_places = df_lang["place_name"].value_counts().nlargest(10).reset_index()
 top_places.columns = ["place_name", "n_reviews"]
 plt.figure(figsize=(9,5))
 sns.barplot(data=top_places, y="place_name", x="n_reviews", palette="viridis")
```

```
plt.title(f"Top 10 Places by Reviews ({label})")
plt.xlabel("# Reviews"); plt.ylabel("")
plt.tight_layout(); plt.show()

def plot_wordcloud(df_lang, label):
 if not HAVE_WORDCLOUD:
 print("Wordcloud not available: package not installed.")
 return
 if df_lang.empty:
 print(f"No {label} reviews for wordcloud.")
 return
 if label.lower().startswith("portuguese"):
 # Rebuild tokens without stemming for Portuguese wordcloud
 tokens = []
 for txt in df_lang["review_text_clean"].dropna():
 toks = simple_tokenize(txt)
 toks = [t for t in toks if t.isalpha() and t not in stop_
 tokens.extend(toks)
 txt = " ".join(tokens)
 else:
 txt = " ".join(df_lang["text_processed"].dropna().tolist())
 if len(txt.strip()) == 0:
 print(f"No text available for wordcloud ({label}).")
 return
 wc = WordCloud(width=1200, height=600, background_color='white',
 plt.figure(figsize=(12,6))
 plt.imshow(wc, interpolation='bilinear')
 plt.axis('off')
 plt.title(f"WordCloud ({label})")
 plt.show()



=== 1. Rating Distributions ===
print("\n==== Rating Distributions ====")
plot_ratings(df_en, "English")
plot_ratings(df_pt, "Portuguese")

=== 2. Sentiment Distributions (Bilingual) ===
print("\n==== Sentiment Distributions ====")
plot_sentiment(df_en, "English", "VADER")
plot_sentiment(df_pt, "Portuguese", "BERTweet-PT")


=== 3. Top Places by Review Count ===
print("\n==== Top Places by Review Count ====")
plot_top_places(df_en, "English")
plot_top_places(df_pt, "Portuguese")


=== 4. Word Clouds (Language-Specific) ===
print("\n==== Word Clouds ====")
plot_wordcloud(df_en, "English")
plot_wordcloud(df_pt, "Portuguese")
```

==== Rating Distributions ====



#### ==== Sentiment Distributions ====






==== Top Places by Review Count ====

/tmp/ipykernel\_5558/2347132612.py:39: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.


```
sns.barplot(data=top_places, y="place_name", x="n_reviews", palette="viridis")
```



/tmp/ipykernel\_5558/2347132612.py:39: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(data=top_places, y="place_name", x="n_reviews", palette="viridis")
```



==== Word Clouds ====



## 6. Save Artifacts & Pipeline Summary

**Purpose:** Persist all pipeline outputs to disk and generate a summary report for reproducibility and downstream analysis.

### Output Files Generated

| File                        | Location        | Contains                                                                                                                                                                                     | Purpose                                                                       |
|-----------------------------|-----------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------|
| <b>reviews_raw.csv</b>      | .../<br>output/ | All extracted reviews with:<br>place_id, place_name, rating,<br>review_text (original),<br>publish_time, lat, lon, place_rating,<br>place_primary_type, poi metadata<br>(gid, amenity, etc.) | Raw API output;<br>enables replay/<br>replay<br>debugging                     |
| <b>reviews_clean.csv</b>    | .../<br>output/ | raw + cleaned: review_text_clean,<br>detected lang, tokens list,<br>token_count, text_processed<br>(stopwords removed, normalized)                                                           | Intermediate<br>format; allows<br>text inspection<br>and validation           |
| <b>reviews_enriched.csv</b> | .../<br>output/ | clean + features: tf_idf features,<br>sentiment_compound (EN via<br>VADER; PT via HuggingFace<br>bertweet-pt-sentiment or TextBlob<br>fallback)                                              | Final analytical<br>dataset; ready<br>for downstream<br>modeling/<br>analysis |

### Downstream Usage

- **reviews\_raw.csv:** Audit API responses; validate geometry/location parsing
- **reviews\_clean.csv:** EDA, text quality checks, language distribution analysis
- **reviews\_enriched.csv:** Train sentiment models, topic analysis

```
In [12]: summary = {
 "raw_reviews_csv": RAW_REVIEWS_CSV if os.path.exists(RAW_REVIEWS_CSV)
 "clean_reviews_csv": CLEAN_REVIEWS_CSV if os.path.exists(CLEAN_REVIEWS_CSV)
 "enriched_reviews_csv": os.path.join(OUTPUT_DIR, "reviews_enriched.csv")
 "n_reviews_raw": int(pd.read_csv(RAW_REVIEWS_CSV).shape[0]) if os.path.exists(RAW_REVIEWS_CSV)
 "n_reviews_clean": int(pd.read_csv(CLEAN_REVIEWS_CSV).shape[0]) if os.path.exists(CLEAN_REVIEWS_CSV)
}
print(json.dumps(summary, indent=2))

{
 "raw_reviews_csv": ".../output/reviews_raw.csv",
 "clean_reviews_csv": ".../output/reviews_clean.csv",
 "enriched_reviews_csv": ".../output/reviews_enriched.csv",
 "n_reviews_raw": 6800,
 "n_reviews_clean": 5298
}
```

## 7. Topic Analysis (LDA)

**Purpose:** Discover latent themes and topics within review texts using Latent Dirichlet Allocation (LDA), enabling automatic identification of recurring discussion themes across different languages.

## Process

### 1. Prepare Corpus:

- Split reviews by language (English and Portuguese separately)
- Use `text_processed` (stopword-removed, normalized text) as input
- Build dictionary of unique tokens per language

### 2. LDA Model Training:

- Fit separate LDA models for English and Portuguese reviews
- Number of topics: 5 (configurable)
- LDA hyperparameters:
  - `passes=10` : Number of iterations over corpus
  - `workers=4` : Parallel processing threads
  - `per_word_topics=True` : Detailed per-word topic assignments

### 3. Topic Extraction:

- Extract top N words per topic (highest probability)
- Compute topic distribution per document (review)
- Identify dominant topic for each review

### 4. Visualization & Analysis:

- Display top words per topic with probabilities
- Show topic proportions across dataset
- Create visualizations of topic coherence and prevalence

## Key Outputs

- **Per-language LDA models:** Trained topic models for EN and PT
- **Topic interpretations:** Most probable words defining each topic
- **Document-topic mappings:** Which topics appear in each review
- **Topic statistics:** Prevalence and coherence metrics

```
In []: if not df_enriched.empty and HAVE_GENSIM:
 print("== Topic Analysis (LDA) ==\n")

 # Configuration
 NUM_TOPICS = 5
 PASSES = 10
 WORKERS = 4
 MIN_WORD_FREQ = 2

 def train_lda_model(df_lang, lang_name, num_topics=NUM_TOPICS):
 """Train LDA model for a given language."""
 if df_lang.empty or len(df_lang) < 5:
 print(f"Insufficient {lang_name} reviews ({len(df_lang)}) for")
 return None, None

 # Convert text_processed to token lists
 texts = [text.split() for text in df_lang["text_processed"].fillna("")]
 texts = [[token for token in text if token] for text in texts] #

 if not texts or all(len(t) == 0 for t in texts):
```

```
 print(f"No tokens available for {lang_name}.")
 return None, None

 # Build dictionary and corpus
 dictionary = corpora.Dictionary(texts)
 dictionary.filter_extremes(no_below=MIN_WORD_FREQ, no_above=0.7,
 corpus = [dictionary.doc2bow(text) for text in texts]

 if not corpus:
 print(f"Empty corpus for {lang_name}.")
 return None, None

 # Train LDA model
 print(f"Training LDA model for {lang_name} ({len(texts)} reviews)
 try:
 lda_model = LdaMulticore(
 corpus=corpus,
 id2word=dictionary,
 num_topics=num_topics,
 random_state=42,
 passes=PASSES,
 workers=WORKERS,
 per_word_topics=True,
 minimum_probability=0.0,
)
 print(f"✓ {lang_name} LDA model trained successfully.\n")
 return lda_model, dictionary
 except Exception as e:
 print(f"x Error training {lang_name} LDA model: {e}\n")
 return None, None

 # Train models for each language
 lda_en, dict_en = train_lda_model(df_en, "English", NUM_TOPICS)
 lda_pt, dict_pt = train_lda_model(df_pt, "Portuguese", NUM_TOPICS)

 # Display topics
 def display_topics(lda_model, lang_name):
 """Print top words for each topic."""
 if lda_model is None:
 print(f"No {lang_name} LDA model available.")
 return

 print(f"\n{'='*60}")
 print(f" {lang_name.upper()} - TOP TOPICS")
 print(f"{'='*60}")
 topics = lda_model.print_topics(num_words=10)
 for topic_id, topic_words in topics:
 print(f"\nTopic {topic_id}:")
 # Parse and display words with weights
 words_weights = [item.split('*') for item in topic_words.split()]
 for weight, word in words_weights:
 word_clean = word.strip('\'')
 print(f" {word_clean:20s} {float(weight):.4f}")

 # Compute dominant topics for each review
 def get_dominant_topics(df_lang, lda_model, dictionary, lang_name):
 """Extract dominant topic per review."""
 if lda_model is None or df_lang.empty:
 return pd.DataFrame()
```

```
texts = [text.split() for text in df_lang["text_processed"].fillna(" ")]
texts = [[token for token in text if token] for text in texts]

dominant_topics = []
for idx, text in enumerate(texts):
 bow = dictionary.doc2bow(text)
 if bow:
 topics = lda_model.get_document_topics(bow)
 if topics:
 dominant_topic = max(topics, key=lambda x: x[1])
 dominant_topics.append({
 "topic_id": dominant_topic[0],
 "topic_prob": dominant_topic[1],
 })
 else:
 dominant_topics.append({"topic_id": None, "topic_prob": 0})
 else:
 dominant_topics.append({"topic_id": None, "topic_prob": 0})

return pd.DataFrame(dominant_topics)

Visualize topic distributions
def plot_topic_distribution(df_topics, lang_name):
 """Plot topic prevalence."""
 if df_topics.empty or df_topics["topic_id"].isna().all():
 print(f"No topics to plot for {lang_name}.")
 return

 topic_counts = df_topics["topic_id"].value_counts().sort_index()
 plt.figure(figsize=(9, 5))
 plt.bar(topic_counts.index, topic_counts.values, color="mediumpurple")
 plt.title(f"Topic Distribution ({lang_name})")
 plt.xlabel("Topic ID")
 plt.ylabel("Number of Reviews")
 plt.xticks(topic_counts.index)
 plt.tight_layout()
 plt.show()

=== Display English Topics ===
display_topics(lda_en, "English")

=== Display Portuguese Topics ===
display_topics(lda_pt, "Portuguese")

=== Get and display dominant topics for English ===
if lda_en is not None:
 topics_en = get_dominant_topics(df_en, lda_en, dict_en, "English")
 if not topics_en.empty:
 print("\nEnglish Topic Distribution:")
 print(topics_en["topic_id"].value_counts().sort_index())
 plot_topic_distribution(topics_en, "English")

=== Get and display dominant topics for Portuguese ===
if lda_pt is not None:
 topics_pt = get_dominant_topics(df_pt, lda_pt, dict_pt, "Portuguese")
 if not topics_pt.empty:
 print("\nPortuguese Topic Distribution:")
 print(topics_pt["topic_id"].value_counts().sort_index())
 plot_topic_distribution(topics_pt, "Portuguese")
```

```
elif not df_enriched.empty and not HAVE_GENSIM:
 print("Topic Analysis skipped: gensim not installed. Install with: pi
else:
 print("No enriched data available for topic analysis.")
```

**==== Topic Analysis (LDA) ====**

Training LDA model for English (2492 reviews)...

✓ English LDA model trained successfully.

Training LDA model for Portuguese (2805 reviews)...

✓ English LDA model trained successfully.

Training LDA model for Portuguese (2805 reviews)...

✓ Portuguese LDA model trained successfully.

---

**ENGLISH - TOP TOPICS**

---

**Topic 0:**

|            |        |
|------------|--------|
| food       | 0.0290 |
| great      | 0.0220 |
| service    | 0.0220 |
| place      | 0.0160 |
| experience | 0.0140 |
| friendly   | 0.0140 |
| staff      | 0.0140 |
| amazing    | 0.0140 |
| good       | 0.0110 |
| pizza      | 0.0100 |

**Topic 1:**

|            |        |
|------------|--------|
| best       | 0.0310 |
| place      | 0.0280 |
| aveiro     | 0.0260 |
| top        | 0.0240 |
| time       | 0.0110 |
| one        | 0.0110 |
| visit      | 0.0110 |
| love       | 0.0100 |
| portuguese | 0.0100 |
| portugal   | 0.0090 |

**Topic 2:**

|           |        |
|-----------|--------|
| good      | 0.0440 |
| food      | 0.0250 |
| service   | 0.0240 |
| great     | 0.0240 |
| nice      | 0.0220 |
| place     | 0.0210 |
| price     | 0.0200 |
| staff     | 0.0180 |
| friendly  | 0.0170 |
| delicious | 0.0130 |

**Topic 3:**

|              |        |
|--------------|--------|
| service      | 0.0170 |
| one          | 0.0120 |
| time         | 0.0120 |
| recommend    | 0.0110 |
| even         | 0.0100 |
| professional | 0.0090 |
| experience   | 0.0080 |

|        |        |
|--------|--------|
| good   | 0.0080 |
| minute | 0.0080 |
| like   | 0.0080 |

## Topic 4:

|           |        |
|-----------|--------|
| room      | 0.0210 |
| good      | 0.0160 |
| clean     | 0.0150 |
| location  | 0.0140 |
| nice      | 0.0120 |
| place     | 0.0120 |
| aveiro    | 0.0110 |
| well      | 0.0110 |
| great     | 0.0100 |
| breakfast | 0.0100 |

=====  
PORTUGUESE - TOP TOPICS  
=====

## Topic 0:

|       |        |
|-------|--------|
| faz   | 0.0130 |
| tod   | 0.0120 |
| dia   | 0.0120 |
| vez   | 0.0090 |
| porqu | 0.0090 |
| outr  | 0.0090 |
| atend | 0.0080 |
| ped   | 0.0080 |
| pod   | 0.0070 |
| diss  | 0.0070 |

## Topic 1:

|        |        |
|--------|--------|
| tod    | 0.0160 |
| compr  | 0.0120 |
| aveir  | 0.0110 |
| cas    | 0.0100 |
| melhor | 0.0090 |
| dia    | 0.0090 |
| bem    | 0.0090 |
| sempr  | 0.0090 |
| client | 0.0080 |
| faz    | 0.0070 |

## Topic 2:

|        |        |
|--------|--------|
| atend  | 0.0530 |
| bom    | 0.0390 |
| excel  | 0.0360 |
| recom  | 0.0250 |
| qual   | 0.0230 |
| serviç | 0.0230 |
| preç   | 0.0200 |
| simpá  | 0.0190 |
| boa    | 0.0170 |
| espaç  | 0.0150 |

## Topic 3:

|        |        |
|--------|--------|
| atend  | 0.0260 |
| client | 0.0140 |
| pesso  | 0.0130 |

|      |        |
|------|--------|
| faz  | 0.0130 |
| hor  | 0.0100 |
| ter  | 0.0100 |
| empr | 0.0090 |
| aind | 0.0090 |
| loj  | 0.0080 |
| cas  | 0.0080 |

**Topic 4:**

|         |        |
|---------|--------|
| profiss | 0.0390 |
| recom   | 0.0280 |
| excel   | 0.0250 |
| sup     | 0.0210 |
| atend   | 0.0200 |
| sempr   | 0.0190 |
| tod     | 0.0160 |
| trabalh | 0.0130 |
| atenci  | 0.0120 |
| melhor  | 0.0120 |

✓ Portuguese LDA model trained successfully.

---

---

ENGLISH - TOP TOPICS

---

---

**Topic 0:**

|            |        |
|------------|--------|
| food       | 0.0290 |
| great      | 0.0220 |
| service    | 0.0220 |
| place      | 0.0160 |
| experience | 0.0140 |
| friendly   | 0.0140 |
| staff      | 0.0140 |
| amazing    | 0.0140 |
| good       | 0.0110 |
| pizza      | 0.0100 |

**Topic 1:**

|            |        |
|------------|--------|
| best       | 0.0310 |
| place      | 0.0280 |
| aveiro     | 0.0260 |
| top        | 0.0240 |
| time       | 0.0110 |
| one        | 0.0110 |
| visit      | 0.0110 |
| love       | 0.0100 |
| portuguese | 0.0100 |
| portugal   | 0.0090 |

**Topic 2:**

|          |        |
|----------|--------|
| good     | 0.0440 |
| food     | 0.0250 |
| service  | 0.0240 |
| great    | 0.0240 |
| nice     | 0.0220 |
| place    | 0.0210 |
| price    | 0.0200 |
| staff    | 0.0180 |
| friendly | 0.0170 |

delicious 0.0130

Topic 3:

|              |        |
|--------------|--------|
| service      | 0.0170 |
| one          | 0.0120 |
| time         | 0.0120 |
| recommend    | 0.0110 |
| even         | 0.0100 |
| professional | 0.0090 |
| experience   | 0.0080 |
| good         | 0.0080 |
| minute       | 0.0080 |
| like         | 0.0080 |

Topic 4:

|           |        |
|-----------|--------|
| room      | 0.0210 |
| good      | 0.0160 |
| clean     | 0.0150 |
| location  | 0.0140 |
| nice      | 0.0120 |
| place     | 0.0120 |
| aveiro    | 0.0110 |
| well      | 0.0110 |
| great     | 0.0100 |
| breakfast | 0.0100 |

=====

===== PORTUGUESE - TOP TOPICS =====

=====

Topic 0:

|       |        |
|-------|--------|
| faz   | 0.0130 |
| tod   | 0.0120 |
| dia   | 0.0120 |
| vez   | 0.0090 |
| porqu | 0.0090 |
| outr  | 0.0090 |
| atend | 0.0080 |
| ped   | 0.0080 |
| pod   | 0.0070 |
| diss  | 0.0070 |

Topic 1:

|        |        |
|--------|--------|
| tod    | 0.0160 |
| compr  | 0.0120 |
| aveir  | 0.0110 |
| cas    | 0.0100 |
| melhor | 0.0090 |
| dia    | 0.0090 |
| bem    | 0.0090 |
| sempr  | 0.0090 |
| client | 0.0080 |
| faz    | 0.0070 |

Topic 2:

|       |        |
|-------|--------|
| atend | 0.0530 |
| bom   | 0.0390 |
| excel | 0.0360 |
| recom | 0.0250 |
| qual  | 0.0230 |

|        |        |
|--------|--------|
| serviç | 0.0230 |
| preç   | 0.0200 |
| simpá  | 0.0190 |
| boa    | 0.0170 |
| espaç  | 0.0150 |

Topic 3:

|        |        |
|--------|--------|
| atend  | 0.0260 |
| client | 0.0140 |
| pesso  | 0.0130 |
| faz    | 0.0130 |
| hor    | 0.0100 |
| ter    | 0.0100 |
| empr   | 0.0090 |
| aind   | 0.0090 |
| loj    | 0.0080 |
| cas    | 0.0080 |

Topic 4:

|         |        |
|---------|--------|
| profiss | 0.0390 |
| recom   | 0.0280 |
| excel   | 0.0250 |
| sup     | 0.0210 |
| atend   | 0.0200 |
| sempr   | 0.0190 |
| tod     | 0.0160 |
| trabalh | 0.0130 |
| atenci  | 0.0120 |
| melhor  | 0.0120 |

English Topic Distribution:


| topic_id | count |
|----------|-------|
| 0.0      | 524   |
| 1.0      | 329   |
| 2.0      | 803   |
| 3.0      | 438   |
| 4.0      | 391   |

Name: count, dtype: int64

English Topic Distribution:

| topic_id | count |
|----------|-------|
| 0.0      | 524   |
| 1.0      | 329   |
| 2.0      | 803   |
| 3.0      | 438   |
| 4.0      | 391   |

Name: count, dtype: int64




Portuguese Topic Distribution:

topic\_id

0.0 335  
1.0 341  
2.0 978  
3.0 336  
4.0 805

Name: count, dtype: int64



English model topics: 5

Portuguese model topics: 5